

Hedge detection as a lens on framing in the GMO debates: A position paper

Eunsol Choi*, Chenhao Tan*, Lillian Lee*, Cristian Danescu-Niculescu-Mizil* and Jennifer Spindel†

*Department of Computer Science, †Department of Plant Breeding and Genetics

Cornell University

ec472@cornell.edu, chenhao|llee|cristian@cs.cornell.edu, jes462@cornell.edu

Abstract

Understanding the ways in which participants in public discussions frame their arguments is important in understanding how public opinion is formed. In this paper, we adopt the position that it is time for more computationally-oriented research on problems involving framing. In the interests of furthering that goal, we propose the following specific, interesting and, we believe, relatively accessible question: In the controversy regarding the use of genetically-modified organisms (GMOs) in agriculture, do pro- and anti-GMO articles differ in whether they choose to adopt a more “scientific” tone?

Prior work on the rhetoric and sociology of science suggests that *hedging* may distinguish popular-science text from text written by professional scientists for their colleagues. We propose a detailed approach to studying whether hedge detection can be used to understand scientific framing in the GMO debates, and provide corpora to facilitate this study. Some of our preliminary analyses suggest that hedges occur less frequently in scientific discourse than in popular text, a finding that contradicts prior assertions in the literature. We hope that our initial work and data will encourage others to pursue this promising line of inquiry.

Publication venue: ACL Workshop on Extra-Propositional Aspects of Meaning in Computational Linguistics, 2012

1 Introduction

1.1 Framing, “scientific discourse”, and GMOs in the media

The issue of *framing* (Goffman, 1974; Scheufele, 1999; Benford and Snow, 2000) is of great importance in understanding how public opinion is formed. In their *Annual Review of Political Science* survey, Chong and Druckman (2007) describe framing effects as occurring “when (often small) changes in the presentation of an issue or an event produce (sometimes large) changes of opinion” (p. 104); as an example, they cite a study wherein respondents answered differently, when asked whether a hate group should be allowed to hold a rally, depending on whether the question was phrased as one of “free speech” or one of “risk of violence”.

The genesis of our work is in a framing question motivated by a relatively current political issue. In media coverage of transgenic crops and the use of genetically modified organisms (GMOs) in food, do pro-GMO vs. anti-GMO articles differ not just with respect to word choice, but in adopting a more “scientific” discourse, meaning the inclusion of more uncertainty and fewer emotionally-laden words? We view this as an interesting question from a text analysis perspective (with potential applications and implications that lie outside the scope of this article).

1.2 Hedging as a sign of scientific discourse

To obtain a computationally manageable characterization of “scientific discourse”, we turned to studies of the culture and language of science, a body of work spanning fields ranging from sociology to

applied linguistics to rhetoric and communication (Gilbert and Mulkay, 1984; Latour, 1987; Latour and Woolgar, 1979; Halliday and Martin, 1993; Bazerman, 1988; Fahnestock, 2004; Gross, 1990).

One characteristic that has drawn quite a bit of attention in such studies is *hedging* (Myers, 1989; Hyland, 1998; Lewin, 1998; Salager-Meyer, 2011).¹ Hyland (1998, pg. 1) defines hedging as the expression of “tentativeness and possibility” in communication, or, to put it another way, language corresponding to “the writer withholding full commitment to statements” (pg. 3). He supplies many real-life examples from scientific research articles, including the following:

1. ‘*It seems that this group plays a critical role in orienting the carboxyl function*’ (emphasis Hyland’s)
2. ‘*...implies that phytochrome A is also not necessary for normal photomorphogenesis, at least under these irradiation conditions*’ (emphasis Hyland’s)
3. ‘*We wish to suggest a structure for the salt of deoxyribose nucleic acid (D.N.A.)*’ (emphasis added)²

Several scholars have asserted the centrality of hedging in scientific and academic discourse, which corresponds nicely to the notion of “more uncertainty” mentioned above. Hyland (1998, p. 6) writes, “Despite a widely held belief that professional scientific writing is a series of impersonal statements of fact which add up to the truth, hedges are abundant in science and play a critical role in academic writing”. Indeed, Myers (1989, p. 13) claims that in scientific research articles, “The hedging of claims is so common that a sentence that looks like a claim but has no hedging is probably not a statement of new knowledge”.³

Not only is understanding hedges important to understanding the rhetoric and sociology of science,

¹In linguistics, hedging has been studied since the 1970s (Lakoff, 1973).

²This example originates from Watson and Crick’s landmark 1953 paper. Although the sentence is overtly tentative, did Watson and Crick truly intend to be polite and modest in their claims? See Varttala (2001) for a review of arguments regarding this question.

³Note the inclusion of the hedge “probably”.

but hedge detection and analysis — in the sense of identifying uncertain or uncertainly-sourced information (Farkas et al., 2010) — has important applications to information extraction, broadly construed, and has thus become an active sub-area of natural-language processing. For example, the CoNLL 2010 Shared Task was devoted to this problem (Farkas et al., 2010).

Putting these two lines of research together, we see before us what appears to be an interesting interdisciplinary and, at least in principle, straightforward research program: relying on the aforementioned rhetoric analyses to presume that hedging is a key characteristic of scientific discourse, build a hedge-detection system to computationally ascertain which proponents in the GMO debate tend to use more hedges and thus, by presumption, tend to adopt a more “scientific” frame.⁴

1.3 Contributions

Our overarching goal in this paper is to convince more researchers in NLP and computational linguistics to work on problems involving framing. We try to do so by proposing a specific problem that may be relatively accessible. Despite the apparent difficulty in addressing such questions, we believe that progress can be made by drawing on observations drawn from previous literature across many fields, and integrating such work with movements in the computational community toward consideration of extra-propositional and pragmatic concerns. We have thus intentionally tried to “cover a lot of ground”, as one referee put it, in the introductory material just discussed.

Since framing problems are indeed difficult, we elected to narrow our scope in the hope of making some partial progress. Our technical goal here, at this workshop, where hedge detection is one of the most relevant topics to the broad questions we have raised, is *not* to learn to classify texts as being pro- vs. anti-GMO, or as being scientific or not, per se.⁵

⁴However, this presumption that more hedges characterize a more scientific discourse has been contested. See section 2 for discussion and section 4.2 for our empirical investigation.

⁵Several other groups have addressed the problem of trying to identify different sides or perspectives (Lin et al., 2006; Hardisty et al., 2010; Beigman Klebanov et al., 2010; Ahmed and Xing, 2010).

Our focus is on whether hedging specifically, considered as a single feature, is correlated with these different document classes, because of the previous research attention that has been devoted to hedging in particular and because of hedging being one of the topics of this workshop. The point of this paper is thus not to compare the efficacy of hedging features with other types, such as bag-of-words features. Of course, to do so is an important and interesting direction for future work.

In the end, we were not able to achieve satisfactory results even with respect to our narrowed goal. However, we believe that other researchers may be able to follow the plan of attack we outline below, and perhaps use the data we are releasing, in order to achieve our goal. We would welcome hearing the results of other people's efforts.

2 How should we test whether hedging distinguishes scientific text?

One very important point that we have not yet addressed is: While the literature agrees on the importance of hedging in scientific text, the *relative degree* of hedging in scientific vs. non-scientific text is a matter of debate.

On the one side, we have assertions like those of Fahnestock (1986), who shows in a clever, albeit small-scale, study involving parallel texts that when scientific observations pass into popular accounts, changes include “removing hedges ... thus conferring greater certainty on the reported facts” (pg. 275). Similarly, Juanillo, Jr. (2001) refers to a shift from a forensic style to a “celebratory” style when scientific research becomes publicized, and credits Brown (1998) with noting that “celebratory scientific discourses tend to pay less attention to caveats, contradictory evidence, and qualifications that are highlighted in forensic or empiricist discourses. By downplaying scientific uncertainty, it [sic] alludes to greater certainty of scientific results for public consumption” (Juanillo, Jr., 2001, p. 42).

However, others have contested claims that the popularization process involves simplification, distortion, hype, and dumbing down, as Myers (2003) colorfully puts it; he provides a critique of the relevant literature. Varttala (1999) ran a corpus analysis in which hedging was found not just in pro-

fessional medical articles, but was also “typical of popular scientific articles dealing with similar topics” (p. 195). Moreover, significant variation in use of hedging has been found across disciplines and authors' native language; see Salager-Meyer (2011) or Varttala (2001) for a review.

To the best of our knowledge, there have been no large-scale empirical studies validating the hypothesis that hedges appear more or less frequently in scientific discourse.

Proposed procedure Given the above, our **first step** must be to determine whether hedges are more or less prominent in “professional scientific” (henceforth “*prof-science*”) vs. “public science” (henceforth “*pop-science*”) discussions of GMOs. Of course, for a large-scale study, finding hedges requires developing and training an effective hedge detection algorithm.

If the first step shows that hedges can indeed be used to effectively distinguish *prof-science* vs. *pop-science* discourse on GMOs, then the **second step** is to examine whether the use of hedging in *pro-GMO* articles follows our inferred “scientific” occurrence patterns to a greater extent than the hedging in *anti-GMO* articles.

However, as our hedge classifier trained on the CoNLL dataset did not perform reliably on the different domain of *prof-science* vs. *pop-science* discussions of GMOs, we focus the main content of this paper on the first step. We describe data collection for the second step in the appendix.

3 Data

To accomplish the first step of our proposed procedure outlined above, we first constructed a *prof-science*/*pop-science* corpus by pulling text from Web of Science for *prof-science* examples and from LexisNexis for *pop-science* examples, as described in Section 3.1. Our corpus will be posted online at <https://confluence.cornell.edu/display/llresearch/HedgingFramingGMOs>.

As noted above, computing the degree of hedging in the aforementioned corpus requires access to a hedge-detection algorithm. We took a supervised approach, taking advantage of the availability of the CoNLL 2010 hedge-detection training and evaluation corpora, described in Section 3.2

Dataset	Doc type	# docs	# sentences	Avg sentence length	Flesch reading ease
Prof-science/pop-science corpus					
WOS	abstracts	648	5596	22.35	23.39
LEXIS	(short) articles	928	36795	24.92	45.78
Hedge-detection corpora					
Bio (train)	abstracts, articles	1273, 9	14541 (18% uncertain)	29.97	20.77
Bio (eval)	articles	15	5003 (16% uncertain)	31.30	30.49
Wiki (train)	paragraphs	2186	11111 (22% uncertain)	23.07	35.23
Wiki (eval)	paragraphs	2346	9634 (23% uncertain)	20.82	31.71

Table 1: Basic descriptive statistics for the main corpora we worked with. We created the first two. Higher Flesch scores indicate text that is easier to read.

3.1 Prof-science/pop-science data: LEXIS and WOS

As mentioned previously, a corpus of prof-science and pop-science articles is required to ascertain whether hedges are more prevalent in one or the other of these two writing styles. Since our ultimate goal is to look at discourse related to GMOs, we restrict our attention to documents on this topic.

Thomson Reuter’s Web of Science (WOS), a database of scientific journal and conference articles, was used as a source of prof-science samples. We chose to collect abstracts, rather than full scientific articles, because intuition suggests that the language in abstracts is more high-level than that in the bodies of papers, and thus more similar to the language one would see in a public debate on GMOs. To select for on-topic abstracts, we used the phrase “transgenic foods” as a search keyword and discarded results containing any of a hand-selected list of off-topic filtering terms (e.g., “mice” or “rats”). We then made use of domain expertise to manually remove off-topic texts. The process yielded 648 documents for a total of 5596 sentences.

Our source of pop-science articles was Lexis-Nexis (LEXIS). On-topic documents were collected from US newspapers using the search keywords “genetically modified foods” or “transgenic crops” and then imposing the additional requirement that at least two terms on a hand-selected list⁷ be present in each document. After the removal of duplicates

and texts containing more than 2000 words to delete excessively long articles, our final pop-science sub-corpus was composed of 928 documents.

3.2 CoNLL hedge-detection training data⁸

As described in Farkas et al. (2010), the motivation behind the CoNLL 2010 shared task is that “distinguishing factual and uncertain information in texts is of essential importance in information extraction”. As “uncertainty detection is extremely important for biomedical information extraction”, one component of the dataset is biological abstracts and full articles from the BioScope corpus (Bio). Meanwhile, the chief editors of Wikipedia have drawn the attention of the public to specific markers of uncertainty known as weasel words⁹: they are words or phrases “aimed at creating an impression that something specific and meaningful has been said”, when, in fact, “only a vague or ambiguous claim, or even a refutation, has been communicated”. An example is “It has been claimed that ...”: the claimant has not been identified, so the source of the claim cannot be verified. Thus, another part of the dataset is a set of Wikipedia articles (Wiki) annotated with weasel-word information. We view the combined Bio+Wiki corpus (henceforth the CoNLL dataset) as valuable for developing hedge detectors, and we attempt to study whether classifiers trained on this data can be generalized to other datasets.

3.3 Comparison

Table 1 gives the basic statistics on the main datasets we worked with. Though WOS and LEXIS differ in

⁷The term list: GMO, GM, GE, genetically modified, genetic modification, modified, modification, genetic engineering, engineered, bioengineered, franken, transgenic, spliced, G.M.O., tweaked, manipulated, engineering, pharming, aquaculture.

⁸<http://www.inf.u-szeged.hu/rgai/conll2010st/>

⁹http://en.wikipedia.org/wiki/Weasel_word

the total number of sentences, the average sentence length is similar. The average sentence length in Bio is longer than that in Wiki. The articles in WOS are markedly more difficult to read than the articles in LEXIS according to Flesch reading ease (Kincaid et al., 1975).

4 Hedging to distinguish scientific text: Initial annotation

As noted in Section 1, it is not a priori clear whether hedging distinguishes scientific text or that more hedges correspond to a more “scientific” discourse. To get an initial feeling for how frequently hedges occur in WOS and LEXIS, we hand-annotated a sample of sentences from each. In Section 4.1, we explain the annotation policy of the CoNLL 2010 Shared Task and our own annotation method for WOS and LEXIS. After that, we move forward in Section 4.2 to compare the percentage of uncertain sentences in prof-science vs. pop-science text on this small hand-labeled sample, and gain some evidence that there is indeed a difference in hedge occurrence rates, although, perhaps surprisingly, there seem to be more hedges in the *pop-science* texts.

As a side benefit, we subsequently use the hand-labeled sample we produce to investigate the accuracy of an automatic hedge detector in the WOS+LEXIS domain; more on this in Section 5.

4.1 Uncertainty annotation

CoNLL 2010 Shared Task annotation policy As described in Farkas et al. (2010, pg. 4), the data annotation policies for the CoNLL 2010 Shared Task were that “sentences containing at least one cue were considered as uncertain, while sentences with no cues were considered as factual”, where a cue is a linguistic marker that *in context* indicates uncertainty. A straightforward example of a sentence marked “uncertain” in the Shared Task is ‘Mild bladder wall thickening *raises the question of* cystitis.’ The annotated cues are not necessarily general, particularly in Wiki, where some of the marked cues are as specific as ‘*some of schumann’s best choral writing*’, ‘*people of the jewish tradition*’, or ‘*certain leisure or cultural activities*’.

Note that “uncertainty” in the Shared Task definition also encompassed phrasing that “creates an

impression that something important has been said, but what is really communicated is vague, misleading, evasive or ambiguous ... [offering] an opinion without any backup or source”. An example of such a sentence, drawn from Wikipedia and marked “uncertain” in the Shared Task, is ‘Some people claim that this results in a better taste than that of other diet colas (most of which are sweetened with aspartame alone).’; Farkas et al. (2010) write, “The ... sentence does not specify the source of the information, it is just the vague term ‘some people’ that refers to the holder of this opinion”.

Our annotation policy We hand-annotated 200 randomly-sampled sentences, half from WOS and half from LEXIS¹⁰, to gauge the frequency with which hedges occur in each corpus. Two annotators each followed the rules of the CoNLL 2010 Shared Task to label sentences as certain, uncertain, or not a proper sentence.¹¹ The annotators agreed on 153 proper sentences of the 200 sentences (75 from WOS and 78 from LEXIS). Cohen’s Kappa (Fleiss, 1981) was 0.67 on the annotation, which means that the consistency between the two annotators was fair or good. However, there were some interesting cases where the two annotators could not agree. For example, in the sentence ‘Cassava is the staple food of tropical Africa and its production, averaged over 24 countries, has increased more than threefold from 1980 to 2005 ...’, one of the annotators believed that “more than” made the sentence uncertain. These borderline cases indicate that the definition of hedging should be carefully delineated in future studies.

4.2 Percentages of uncertain sentences

To validate the hypothesis that prof-science articles contain more hedges, we computed the percentage of uncertain sentences in our labeled data. As shown in Table 2, we observed a trend contradicting earlier studies. Uncertain sentences were more frequent in LEXIS than in WOS, though the difference was

¹⁰We took steps to attempt to hide from the annotators any explicit clues as to the source of individual sentences: the subset of authors who did the annotation were not those that collected the data, and the annotators were presented the sentences in random order.

¹¹The last label was added because of a few errors in scraping the data.

Dataset	% of uncertain sentences
WOS	(estimated from 75-sentence sample) 20
LEXIS	(estimated from 78-sentence sample) 28
Bio	17
Wiki	23

Table 2: Percentages of uncertain sentences.

not statistically significant¹² (perhaps not surprising given the small sample size). The same trend was seen in the CoNLL dataset: there, too, the percentage of uncertain sentences was significantly smaller in Bio (prof-science articles) than in Wiki. In order to make a stronger argument about prof-science vs pop-science, however, more annotation on the WOS and LEXIS datasets is needed.

5 Experiments

As stated in Section 1, our proposal requires developing an effective hedge detection algorithm. Our approach for the preliminary work described in this paper is to re-implement Georgescu’s (2010) algorithm; the experimental results on the Bio+Wiki domain, given in Section 5.1, are encouraging. Then we use this method to attempt to validate (at a larger scale than in our manual pilot annotation) whether hedges can be used to distinguish between prof-science and pop-science discourse on GMOs. Unfortunately, our results, given in Section 5.2, are inconclusive, since our trained model could not achieve satisfactory automatic hedge-detection accuracy on the WOS+LEXIS domain.

5.1 Method

We adopted the method of Georgescu (2010): Support Vector Machine classification based on a Gaussian Radial Basis kernel function (Vapnik, 1998; Fan et al., 2005), employing n-grams from annotated cue phrases as features, as described in more detail below. This method achieved the top performance in the CoNLL 2010 Wikipedia hedge-detection task (Farkas et al., 2010), and SVMs have been proven effective for many different applications. We used the LIBSVM toolkit in our experiments¹³.

¹²Throughout, “statistical significance” refers to the student t-test with $p < .05$.

¹³<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>

As described in Section 3.2, there are two separate datasets in the CoNLL dataset. We experimented on them separately (Bio, Wiki). Also, to make our classifier more generalizable to different datasets, we also trained models based on the two datasets combined (Bio+Wiki). As for features, we took advantage of the observation in Georgescu (2010) that the bag-of-words model does not work well for this task. We used different sets of features based on hedge *cue words* that have been annotated as part of the CoNLL dataset distribution¹⁴. The basic feature set was the frequency of each hedge cue word from the training corpus after removing stop words and punctuation and transforming words to lowercase. Then, we extracted unigrams, bigrams and trigrams from each hedge cue phrase. Table 3 shows the number of features in different settings. Notice that there are many more features in Wiki. As mentioned above, in Wiki, some cues are as specific as ‘some of schumann’s best choral writing’, ‘people of the jewish tradition’, or ‘certain leisure or cultural activities’. Taking n-grams from such specific cues can cause some sentences to be classified incorrectly.

Feature source	#features
Bio	220
Bio (cues + bigram + trigram)	340
Wiki	3740
Wiki (cues + bigram + trigram)	10603

Table 3: Number of features.

Best cross-validation performance				
Dataset	(C, γ)	P	R	F
Bio	$(40, 2^{-3})$	84.0	92.0	87.8
Wiki	$(30, 2^{-6})$	64.0	76.3	69.6
Bio+Wiki	$(10, 2^{-4})$	66.7	78.3	72.0

Table 4: Best 5-fold cross-validation performance for Bio and/or Wiki after parameter tuning. As a reminder, we repeat that our intended final test set is the WOS+LEXIS corpus, which is disjoint from Bio+Wiki.

We adopted several techniques from Georgescu

¹⁴For the Bio model, we used cues extracted from Bio. Likewise, the Wiki model used cues from Wiki, and the Bio+Wiki model used cues from Bio+Wiki.

Evaluation set	Model	P	R	F
WOS+LEXIS	Bio	54	68	60
WOS+LEXIS	Wiki	38	54	45
WOS+LEXIS	Bio+Wiki	21	93	34
Sub-corpus performance of the model based on Bio				
WOS	Bio	58	73	65
LEXIS	Bio	52	64	57

Table 5: The upper part shows the performance on WOS and LEXIS based on models trained on the CoNLL dataset. The lower part gives the sub-corpus results for Bio, which provided the best performance on the full WOS+LEXIS corpus.

(2010) to optimize performance through cross validation. Specifically, we tried different combinations of feature sets (the cue phrases themselves, cues + unigram, cues + bigram, cues + trigram, cues + unigram + bigram + trigram, cues + bigram + trigram). We tuned the width of the RBF kernel (γ) and the regularization parameter (C) via grid search over the following range of values: $\{2^{-9}, 2^{-8}, 2^{-7}, \dots, 2^4\}$ for γ , $\{1, 10, 20, 30, \dots, 150\}$ for C . We also tried different weighting strategies for negative and positive classes (i.e., either proportional to the number of positive instances, or uniform). We performed 5-fold cross validation for each possible combination of experimental settings on the three datasets (Bio, Wiki, Bio+Wiki).

Table 4 shows the best performance on all three datasets and the corresponding parameters. In the three datasets, cue+bigram+trigram provided the best performance, and the weighted model consistently produced superior results to the uniform model. The F1 measure for Bio was 87.8, which was satisfactory, while the F1 results for Wiki were 69.6, which were the worst of all the datasets. This resonates with our observation that the task on Wikipedia is more subtly defined and thus requires a more sophisticated approach than counting the occurrences of bigrams and trigrams.

5.2 Results on WOS+LEXIS

Next, we evaluated whether our best classifier trained on the CoNLL dataset can be generalized to other datasets, in particular, the WOS and LEXIS corpus. Performance was measured on the 153 sen-

Evaluation set	(C, γ)	P	R	F
WOS + LEXIS	$(50, 2^{-9})$	68	62	65
WOS	$(50, 2^{-9})$	85	73	79
LEXIS	$(50, 2^{-9})$	57	54	56

Table 6: Best performance after parameter tuning based on the 153 labeled WOS+LEXIS sentences; this gives some idea of the upper-bound potential of our Georgescu-based method. The training set is Bio, which gave the best performance in Table 5.

tences on which our annotators agreed, a dataset that was introduced in Section 4.1. Table 5 shows how the best models trained on Bio, Wiki, and Bio+Wiki, respectively, performed on the 153 labeled sentences. First, we can see that the performance degraded significantly compared to the performance for in-domain cross validation. Second, of the three different models, Bio showed the best performance. Bio+Wiki gave the worst performance, which hints that combining two datasets and cue words may not be a promising strategy: although Bio+Wiki shows very good recall, this can be attributed to its larger feature set, which contains all available cues and perhaps as a result has a very high false-positive rate. We further investigated and compared performance on LEXIS and WOS for the best model (Bio). Not surprisingly, our classifier works better in WOS than in LEXIS.

It is clear that there exist domain differences between the CoNLL dataset and WOS+LEXIS. To better understand the poor cross-domain performance of the classifier, we tuned another model based on the performance on the 153 labeled sentences using Bio as training data. As we can see in Table 6, the performance on WOS improved significantly, while the performance on LEXIS decreased. This is probably caused by the fact that WOS is a collection of scientific paper abstracts, which is more similar to the training corpus than LEXIS, which is a collection of news media articles¹⁵. Also, LEXIS articles are hard to classify even with the tuned model, which challenges the effectiveness of a cue-words frequency approach beyond professional scientific texts. Indeed, the simplicity of our reim-

¹⁵The Wiki model performed better on LEXIS than on WOS. Though the performance was not good, this result further reinforces the possibility of a domain-dependence problem.

plementation of Georgescu’s algorithm seems to cause longer sentences to be classified as uncertain, because cue phrases (or n-grams extracted from cue phrases) are more likely to appear in lengthier sentences. Analysis of the best performing model shows that the false-positive sentences are significantly longer than the false-negative ones.¹⁶

Dataset	Model	% classified uncertain
WOS	Bio	16
LEXIS	Bio	19
WOS	Tuned	15
LEXIS	Tuned	14

Table 7: For completeness, we report here the percentage of uncertain sentences in WOS and LEXIS according to our trained classifiers, although we regard these results as unreliable since those classifiers have low accuracy. Bio refers to the best model trained on Bio only in Section 5.1, while Tuned refers to the model in Table 6 that is tuned based on the 153 labeled sentences in WOS+LEXIS.

While the cross-domain results were not reliable, we produced preliminary results on whether there exist fewer hedges in scientific text. We can see that the relative difference in certain/uncertain ratios predicted by the two different models (Bio, Tuned) are different in Table 7. In the tuned model, the difference between LEXIS and WOS in terms of the percentage of uncertain sentences was not statistically significant, while in the Bio model, their difference was statistically significant. Since the performance of our hedge classifier on the 153 hand-annotated WOS+LEXIS sentences was not reliable, though, we must abstain from making conclusive statements here.

6 Conclusion and future work

In this position paper, we advocated that researchers apply hedge detection not only to the classic motivation of information-extraction problems, but also to questions of how public opinion forms. We proposed a particular problem in how participants in debates frame their arguments. Specifically, we asked whether pro-GMO and anti-GMO articles differ in adopting a more “scientific” discourse. Inspired by

earlier studies in social sciences relating hedging to texts aimed at professional scientists, we proposed addressing the question with automatic hedge detection as a first step. To develop a hedge classifier, we took advantage of the CoNLL dataset and a small annotated WOS and LEXIS dataset. Our preliminary results show there may exist a gap which indicates that hedging may, in fact, distinguish prof-science and pop-science documents. In fact, this computational analysis suggests the possibility that hedges occur less frequently in scientific prose, which contradicts several prior assertions in the literature.

To confirm the argument that pop-science tends to use more hedging than prof-science, we need a hedge classifier that performs more reliably in the WOS and LEXIS dataset than ours does. An interesting research direction would be to develop transfer-learning techniques to generalize hedge classifiers for different datasets, or to develop a general hedge classifier relatively robust to domain differences. In either case, more annotated data on WOS and LEXIS is needed for better evaluation or training.

Another strategy would be to bypass the first step, in which we determine whether hedges are more or less prominent in scientific discourse, and proceed directly to labeling and hedge-detection in pro-GMO and anti-GMO texts. However, this will not answer the question of whether advocates in debates other than on GMO-related topics employ a more scientific discourse. Nonetheless, to aid those who wish to pursue this alternate strategy, we have collected two sets of opinionated articles on GMO (pro- and anti-); see appendix for more details.

Acknowledgments We thank Daniel Hopkins and Bonnie Webber for reference suggestions, and the anonymous reviewers for helpful and thoughtful comments. This paper is based upon work supported in part by US NSF grants IIS-0910664 and IIS-1016099, a US NSF graduate fellowship to JS, Google, and Yahoo!

References

Amr Ahmed and Eric P Xing. Staying informed: supervised and semi-supervised multi-view topical

¹⁶Average length of true positive sentences : 28.6, false positive sentences 35.09, false negative sentences: 22.0.

- analysis of ideological perspective. In *EMNLP*, pages 1140–1150, 2010.
- Charles Bazerman. *Shaping Written Knowledge: The Genre and Activity of the Experimental Article in Science*. University of Wisconsin Press, Madison, Wis., 1988.
- Beata Beigman Klebanov, Eyal Beigman, and Daniel Diermeier. Vocabulary choice as an indicator of perspective. In *ACL Short Papers*, pages 253–257, Stroudsburg, PA, USA, 2010. Association for Computational Linguistics.
- Robert D. Benford and David A. Snow. Framing processes and social movements: An overview and assessment. *Annual Review of Sociology*, 26: 611–639, 2000.
- Richard Harvey Brown. *Toward a democratic science: Scientific narration and civic communication*. Yale University Press, New Haven, 1998.
- Dennis Chong and James N. Druckman. Framing theory. *Annual Review of Political Science*, 10: 103–126, 2007.
- Jeanne Fahnestock. Accommodating Science. *Written Communication*, 3(3):275–296, 1986.
- Jeanne Fahnestock. Preserving the figure: Consistency in the presentation of scientific arguments. *Written Communication*, 21(1):6–31, 2004.
- Rong-En Fan, Pai-Hsuen Chen, and Chih-Jen Lin. Working set selection using second order information for training support vector machines. *JMLR*, 6:1889–1918, December 2005. ISSN 1532-4435.
- Richárd Farkas, Veronika Vincze, György Móra, János Csirik, and György Szarvas. The CoNLL-2010 shared task: Learning to detect hedges and their scope in natural language text. In *CoNLL—Shared Task*, pages 1–12, 2010.
- Joseph L. Fleiss. *Statistical Methods for Rates and Proportions*. Wiley series in probability and mathematical statistics. John Wiley & Sons, New York, second edition, 1981.
- Maria Georgescu. A hedgehop over a max-margin framework using hedge cues. In *CoNLL—Shared-Task*, pages 26–31, 2010.
- G. Nigel Gilbert and Michael Joseph Mulkay. *Opening Pandora’s box: A sociological analysis of scientists’ discourse*. CUP Archive, 1984.
- Erving Goffman. *Frame analysis: An essay on the organization of experience*. Harvard University Press, 1974.
- Alan G. Gross. *The rhetoric of science*. Harvard University Press, Cambridge, Mass., 1990.
- Michael Alexander Kirkwood Halliday and James R. Martin. *Writing science: Literacy and discursive power*. Psychology Press, London [u.a.], 1993.
- Eric A Hardisty, Jordan Boyd-Graber, and Philip Resnik. Modeling perspective using adaptor grammars. In *EMNLP*, pages 284–292, 2010.
- Ken Hyland. *Hedging in scientific research articles*. John Benjamins Pub. Co., Amsterdam; Philadelphia, 1998.
- Napoleon K. Juanillo, Jr. Frames for Public Discourse on Biotechnology. In *Genetically Modified Food and the Consumer: Proceedings of the 13th meeting of the National Agricultural Biotechnology Council*, pages 39–50, 2001.
- J. Peter Kincaid, Robert P. Fishburne, Richard L. Rogers, and Brad S. Chissom. Derivation of new readability formulas for navy enlisted personnel. Technical report, National Technical Information Service, Springfield, Virginia, February 1975.
- George Lakoff. Hedges: A study in meaning criteria and the logic of fuzzy concepts. *Journal of Philosophical Logic*, 2(4):458–508, 1973.
- Bruno Latour. *Science in action: How to follow scientists and engineers through society*. Harvard University Press, Cambridge, Mass., 1987.
- Bruno Latour and Steve Woolgar. *Laboratory life: The social construction of scientific facts*. Sage Publications, Beverly Hills, 1979.
- Beverly A. Lewin. Hedging: Form and function in scientific research texts. In *Genre Studies in English for Academic Purposes*, volume 9, pages 89–108. Universitat Jaume I, 1998.
- Wei-Hao Lin, Theresa Wilson, Janyce Wiebe, and Alexander Hauptmann. Which side are you on? identifying perspectives at the document and sentence levels. In *CoNLL*, 2006.

- Greg Myers. The pragmatics of politeness in scientific articles. *Applied Linguistics*, 10(1):1–35, 1989.
- Greg Myers. Discourse studies of scientific popularization: Questioning the boundaries. *Discourse Studies*, 5(2):265–279, 2003.
- Françoise Salager-Meyer. Scientific discourse and contrastive linguistics: hedging. *European Science Editing*, 37(2):35–37, 2011.
- Dietram A. Scheufele. Framing as a theory of media effects. *Journal of Communication*, 49(1):103–122, 1999.
- Vladimir N. Vapnik. *Statistical Learning Theory*. Wiley-Interscience, 1998.
- Teppo Varttala. Remarks on the communicative functions of hedging in popular scientific and specialist research articles on medicine. *English for Specific Purposes*, 18(2):177–200, 1999.
- Teppo Varttala. *Hedging in scientifically oriented discourse: Exploring variation according to discipline and intended audience*. PhD thesis, University of Tampere, 2001.

After the initial collection of data, near-duplicates and irrelevant articles were filtered through clustering, keyword searches and distance between word vectors at the document level. We have collected 762 “anti” documents and 671 “pro” documents. We reduced this to a 404 “pro” and 404 “con” set as follows. Each retained “document” consists of only the first 200 words after excluding the first 50 words of documents containing over 280 words. This was done to avoid irrelevant sections such as *Educators have permission to reprint articles for classroom use; other users, please contact editor@actionbioscience.org for reprint permission. See reprint policy.*

The data will be posted online at <https://confluence.cornell.edu/display/llresearch/HedgingFramingGMOs>.

7 Appendix: pro- vs. anti-GMO dataset

Here, we describe the pro- vs. anti-GMO dataset we collected, in the hopes that this dataset may prove helpful in future research regarding the GMO debates, even though we did not use the corpus in the project described in this paper.

The second step of our overall procedure outlined in the introduction — that step being to examine whether the use of hedging in pro-GMO articles corresponds with our inferred “scientific” occurrence patterns more than that in anti-GMO articles — requires a collection of opinionated articles on GMOs. Our first attempt to use news media articles (LEXIS) was unsatisfying, as we found many articles attempt to maintain a neutral position. This led us to collect documents from more strongly opinionated organizational websites such as Greenpeace (anti-GMO), Non GMO Project (anti-GMO), or Why Biotechnology (pro-GMO). Articles were collected from 20 pro-GMO and 20 anti-GMO organizational web sites.